

# Data statistics for query optimization & Views: maintenance and use


Lecture 6  
2ID35, Spring 2015

George Fletcher

Faculteit Wiskunde & Informatica  
Technische Universiteit Eindhoven

8 May 2015

# The life of a query




# Where we've been

query evaluation

- ▶ physical processing of relational operators

# The life of a query: evaluation



# Where we're headed

query *optimization*

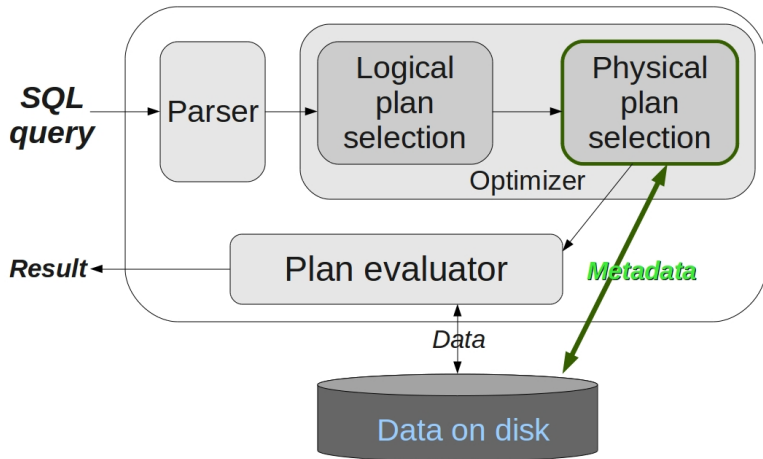
- ▶ logical optimization
- ▶ physical optimization

# Where we're headed

## query *optimization*

- ▶ logical optimization
- ▶ physical optimization
  - ▶ today: cost estimation using data statistics

# The life of a query: cost estimation



# The life of a query: cost estimation

- ▶ A plan is a tree of physical operators
  - ▶ i.e., operators which access and manipulate physical data



# The life of a query: cost estimation

- ▶ A plan is a tree of physical operators
  - ▶ i.e., operators which access and manipulate physical data
- ▶ physical plan indicates
  - ▶ algorithm for each node
  - ▶ the way stored data is obtained (i.e., access paths)
  - ▶ the way in which data is passed between nodes
  - ▶ the order in which nodes are performed

# The life of a query: cost estimation

- ▶ two parts to estimating the cost of a plan
  - ▶ for each node, estimate the cost of performing the operation
  - ▶ for each node, estimate the result size, and any properties it might have (e.g., sorted)

# The life of a query: cost estimation

- ▶ two parts to estimating the cost of a plan
  - ▶ for each node, estimate the cost of performing the operation
  - ▶ for each node, estimate the result size, and any properties it might have (e.g., sorted)
- ▶ the overall estimate is obtained by combining these local estimates
  - ▶ note, however, that errors in estimates propagate exponentially ...

# The life of a query: cost estimation

- ▶ two parts to estimating the cost of a plan
  - ▶ for each node, estimate the cost of performing the operation
  - ▶ for each node, estimate the result size, and any properties it might have (e.g., sorted)
- ▶ the overall estimate is obtained by combining these local estimates
  - ▶ note, however, that errors in estimates propagate exponentially ...
- ▶ keep in mind:
  - ▶ cost estimates are truly approximations
  - ▶ goal is really to just avoid the *worst* plans

# The life of a query: cost estimation

reduction factor:

- ▶ the fraction of tuples which satisfy a condition  $C$  is called the reduction factor of  $C$

# The life of a query: cost estimation

reduction factor:

- ▶ the fraction of tuples which satisfy a condition  $C$  is called the reduction factor of  $C$
- ▶ collectively, the reduction factor of  $C_1 \wedge \dots \wedge C_m$  is  $rf_1 \times \dots \times rf_m$  (assuming statistical independence)

# The life of a query: cost estimation

```
SELECT A1, ..., Ak  
FROM R1, ..., Rn  
WHERE C1 AND ... AND Cm
```

# The life of a query: cost estimation

```
SELECT A1, ..., Ak  
FROM R1, ..., Rn  
WHERE C1 AND ... AND Cm
```

max size:

$$T_{R_1} \times \cdots \times T_{R_n}$$



# The life of a query: cost estimation

```
SELECT A1, ..., Ak  
FROM R1, ..., Rn  
WHERE C1 AND ... AND Cm
```

max size:

$$T_{R_1} \times \cdots \times T_{R_n}$$

estimate of actual size:

$$(rf_1 \times \cdots \times rf_m) \times (T_{R_1} \times \cdots \times T_{R_n})$$

# The life of a query: cost estimation

- ▶ for each RA implementation, we discussed cost

# The life of a query: cost estimation

- ▶ for each RA implementation, we discussed cost
- ▶ the parameters to cost were
  - ▶ input relation size (pages and/or tuples), and
  - ▶ available buffer space

# The life of a query: cost estimation

- ▶ for each RA implementation, we discussed cost
- ▶ the parameters to cost were
  - ▶ input relation size (pages and/or tuples), and
  - ▶ available buffer space
- ▶ let's consider “quick-and-dirty” heuristics for estimating result size (under the assumption of uniform distribution of values)

# Result size estimation: $\pi$

For **projection**, actually computable

# Result size estimation: $\pi$

For **projection**, actually computable

- ▶ reduction factor is  $\frac{\text{new tuple size}}{\text{old tuple size}}$

# Result size estimation: $\pi$

For **projection**, actually computable

- ▶ reduction factor is  $\frac{\text{new tuple size}}{\text{old tuple size}}$
- ▶ impacts the number of pages in output

# Result size estimation: $\sigma$

## selection

- ▶ let  $V(R, A)$  denote the number of distinct values appearing in attribute  $A$  of relation  $R$



# Result size estimation: $\sigma$

## selection

- ▶ let  $V(R, A)$  denote the number of distinct values appearing in attribute  $A$  of relation  $R$ 
  - ▶ by default, we choose  $V(R, A) = 10$ 
    - ▶ some systems keep actual counts
  - ▶ if there is an index on  $R.A$ , then  $V(R, A)$  is equal to the number of keys

# Result size estimation: $\sigma$

## selection

- ▶ let  $V(R, A)$  denote the number of distinct values appearing in attribute  $A$  of relation  $R$ 
  - ▶ by default, we choose  $V(R, A) = 10$ 
    - ▶ some systems keep actual counts
  - ▶ if there is an index on  $R.A$ , then  $V(R, A)$  is equal to the number of keys
- ▶ then, the size estimate of  $\sigma_{A=c}(R)$  is

$$\frac{1}{V(R, A)} T_R$$

# Result size estimation: $\sigma$

For the size estimate of  $\sigma_{A>c}(R)$

- ▶ if  $A$  is not an arithmetic type, then

$$\frac{1}{3} T_R$$

# Result size estimation: $\sigma$

For the size estimate of  $\sigma_{A>c}(R)$

- ▶ if  $A$  is not an arithmetic type, then

$$\frac{1}{3} T_R$$

- ▶ else

$$\frac{\text{highVal}(A) - c}{\text{highVal}(A) - \text{lowVal}(A)} T_R$$

# Result size estimation: $\sigma$

- ▶ for the size estimate of  $\sigma_{A \neq c}(R)$

$$T_R$$

# Result size estimation: $\sigma$

- ▶ for the size estimate of  $\sigma_{A \neq c}(R)$

$$T_R$$

- ▶ for the size estimate of  $\sigma_{C_1 \vee C_2}(R)$

$$\min \{ T_r, \text{size}(\sigma_{C_1}(R)) + \text{size}(\sigma_{C_2}(R)) \}$$

## Result size estimation: $\sigma$

- ▶ for the size estimate of  $\sigma_{A \neq c}(R)$

$$T_R$$

- ▶ for the size estimate of  $\sigma_{C_1 \vee C_2}(R)$

$$\min \{ T_r, \text{size}(\sigma_{C_1}(R)) + \text{size}(\sigma_{C_2}(R)) \}$$

- ▶ for the size estimate of  $\sigma_{C_1 \wedge C_2}(R)$

$$rf_{C_1} \times rf_{C_2} \times T_r$$

## Result size estimation: $\cup$

$$|R \cup S| \approx \frac{(T_R + T_S) + \max\{T_R, T_S\}}{2}$$



## Result size estimation: $\cup$

$$\begin{aligned} |R \cup S| &\approx \frac{(T_R + T_S) + \max\{T_R, T_S\}}{2} \\ &= \frac{(2 \times \text{larger}) + \text{smaller}}{2} \end{aligned}$$

## Result size estimation: $\cup$

$$\begin{aligned} |R \cup S| &\approx \frac{(T_R + T_S) + \max\{T_R, T_S\}}{2} \\ &= \frac{(2 \times \text{larger}) + \text{smaller}}{2} \\ &= \text{larger} + \frac{\text{smaller}}{2} \end{aligned}$$

## Result size estimation: $\cap$

For intersection  $R \cap S$ , the minimum result size is 0, and the max size is  $\min\{T_R, T_S\}$ .

## Result size estimation: $\cap$

For intersection  $R \cap S$ , the minimum result size is 0, and the max size is  $\min\{T_R, T_S\}$ .

Hence,

$$|R \cap S| \approx \frac{0 + \min\{T_R, T_S\}}{2}$$

## Result size estimation: $\cap$

For intersection  $R \cap S$ , the minimum result size is 0, and the max size is  $\min\{T_R, T_S\}$ .

Hence,

$$\begin{aligned} |R \cap S| &\approx \frac{0 + \min\{T_R, T_S\}}{2} \\ &= \frac{\text{smaller}}{2} \end{aligned}$$

## Result size estimation: —

For difference  $R - S$ , the minimum result size is  $T_R - T_S$ , and the max size is  $T_R$ .

## Result size estimation: —

For difference  $R - S$ , the minimum result size is  $T_R - T_S$ , and the max size is  $T_R$ .

Hence,

$$|R - S| \approx \frac{T_R + (T_R - T_S)}{2}$$

## Result size estimation: —

For difference  $R - S$ , the minimum result size is  $T_R - T_S$ , and the max size is  $T_R$ .

Hence,

$$\begin{aligned}|R - S| &\approx \frac{T_R + (T_R - T_S)}{2} \\ &= T_R - \frac{1}{2}T_S\end{aligned}$$



## Result size estimation: $\bowtie$

For natural join  $R \bowtie S$  on attribute  $Y$ :

$$|R \bowtie S| \approx \min \left\{ T_R \frac{T_S}{V(S, Y)}, T_S \frac{T_R}{V(R, Y)} \right\}$$

## Result size estimation: $\bowtie$

For natural join  $R \bowtie S$  on attribute  $Y$ :

$$|R \bowtie S| \approx \min \left\{ T_R \frac{T_S}{V(S, Y)}, T_S \frac{T_R}{V(R, Y)} \right\}$$

since,

- ▶ if  $V(R, Y) = V(S, Y)$ , then each tuple of  $R$  joins with approximately  $\frac{T_S}{V(S, Y)}$  tuples of  $R$ 
  - ▶ and, each tuple of  $S$  joins with approximately  $\frac{T_R}{V(R, Y)}$  tuples of  $S$

## Result size estimation: $\bowtie$

For natural join  $R \bowtie S$  on attribute  $Y$ :

$$|R \bowtie S| \approx \min \left\{ T_R \frac{T_S}{V(S, Y)}, T_S \frac{T_R}{V(R, Y)} \right\}$$

since,

- ▶ if  $V(R, Y) = V(S, Y)$ , then each tuple of  $R$  joins with approximately  $\frac{T_S}{V(S, Y)}$  tuples of  $R$ 
  - ▶ and, each tuple of  $S$  joins with approximately  $\frac{T_R}{V(R, Y)}$  tuples of  $S$
- ▶ and if not, then we minimize the contribution of dangling tuples

# Result size estimation: assumptions made

We have made a few assumptions in these estimates:

# Result size estimation: assumptions made

We have made a few assumptions in these estimates:

- ▶ values across columns are not correlated
  - ▶ this assumption is hard to lift (active area of research)

# Result size estimation: assumptions made

We have made a few assumptions in these estimates:

- ▶ values across columns are not correlated
  - ▶ this assumption is hard to lift (active area of research)
- ▶ values in a single column are uniformly distributed
  - ▶ this assumption can be lifted with **better statistics**

# Result size estimation: histograms

**Histograms:** simple data structures for more refined computation of reduction factors

# Result size estimation: histograms

**Histograms:** simple data structures for more refined computation of reduction factors

Provides approximation of value distribution of an attribute in a relation instance

- ▶ *small*: typically fit on one disk page
- ▶ *accurate*: typically, less than 5% error



# Result size estimation: histograms

**Histograms:** simple data structures for more refined computation of reduction factors

Provides approximation of value distribution of an attribute in a relation instance

- ▶ *small*: typically fit on one disk page
- ▶ *accurate*: typically, less than 5% error

two basic types:

- ▶ equi-width
- ▶ equi-depth

# Result size estimation: histograms

Equi-width, on column  $A$

- ▶ divide range of values appearing in  $A$  into equal sized sub-ranges
- ▶ compute and store total number of tuples falling into each of these “buckets”

# Result size estimation: histograms

Equi-width, on column  $A$

- ▶ divide range of values appearing in  $A$  into equal sized sub-ranges
- ▶ compute and store total number of tuples falling into each of these “buckets”
- ▶ to estimate the output cardinality of a range query on  $A$ , find starting bucket, and scan forward until ending bucket is identified
- ▶ sum number of tuples seen, assuming uniform distribution of values within buckets

# Result size estimation: histograms

Equi-depth, on column  $A$

- ▶ divide range of values appearing in  $A$  into sub-ranges, such that each bucket contains the same number of tuples

# Result size estimation: histograms

Equi-depth, on column  $A$

- ▶ divide range of values appearing in  $A$  into sub-ranges, such that each bucket contains the same number of tuples
- ▶ use the same algorithm to approximate the number of tuples falling in a range query over  $A$

# Result size estimation: histograms

**Example.** Consider a “Sales” attribute with the following actual value distribution:

Number of tuples	Sales value
10	0.5 mil
10	1 mil
10	2 mil
5	5 mil
5	7 mil
5	15 mil
2	40 mil
1	50 mil
1	70 mil
1	100 mil

# Result size estimation: histograms

Suppose we have enough space to store histograms with five buckets

# Result size estimation: histograms

Suppose we have enough space to store histograms with five buckets

equi-width		equi-depth	
value range	tuple count	value range	tuple count
0-20	45	0-0.5	10
20-40	2	0.5-1	10
40-60	1	1-2	10
60-80	1	2-7	10
80-100	1	7-100	10



## Result size estimation: histograms

Suppose we have enough space to store histograms with five buckets

equi-width		equi-depth	
value range	tuple count	value range	tuple count
0-20	45	0-0.5	10
20-40	2	0.5-1	10
40-60	1	1-2	10
60-80	1	2-7	10
80-100	1	7-100	10

The selectivity estimate of “sales  $\leq$  10 million” for equi-width is  $\frac{23}{50} = 46\%$  and for equi-depth is  $\frac{40}{50} = 80\%$

## Result size estimation: histograms

Suppose we have enough space to store histograms with five buckets

equi-width		equi-depth	
value range	tuple count	value range	tuple count
0-20	45	0-0.5	10
20-40	2	0.5-1	10
40-60	1	1-2	10
60-80	1	2-7	10
80-100	1	7-100	10

The selectivity estimate of “sales  $\leq$  10 million” for equi-width is  $\frac{23}{50} = 46\%$  and for equi-depth is  $\frac{40}{50} = 80\%$  which is the actual selectivity!

## Result size estimation: histograms

Suppose we have enough space to store histograms with five buckets

equi-width		equi-depth	
value range	tuple count	value range	tuple count
0-20	45	0-0.5	10
20-40	2	0.5-1	10
40-60	1	1-2	10
60-80	1	2-7	10
80-100	1	7-100	10

The selectivity estimate of “sales  $\leq$  5 million” for equi-width is  $\frac{12}{50} = 24\%$  and for equi-depth is  $\frac{35}{50} = 70\%$

## Result size estimation: histograms

Suppose we have enough space to store histograms with five buckets

equi-width		equi-depth	
value range	tuple count	value range	tuple count
0-20	45	0-0.5	10
20-40	2	0.5-1	10
40-60	1	1-2	10
60-80	1	2-7	10
80-100	1	7-100	10

The selectivity estimate of “sales  $\leq$  5 million” for equi-width is  $\frac{12}{50} = 24\%$  and for equi-depth is  $\frac{35}{50} = 70\%$  which is again the actual selectivity!

# Result size estimation: histograms

## Equi-depth:

- ▶ more accurate than equi-width, since buckets with frequently occurring values correspond to smaller ranges, hence giving finer approximations
- ▶ effective, simple approach to selectivity estimation, and hence quite common

# Result size estimation: histograms

**Exercise.** Consider a “Friends” attribute with the following actual value distribution:

Friends value	Number of tuples
0	1
1	3
2	6
3	10
4	3
5	2
6	2
7	2
8	1

Construct equi-depth and equi-width histograms over this attribute, using three buckets. What estimates do they give for the count of tuples with “friends  $\geq 4$ ”?

# Wrap up

- ▶ Cost estimation, towards choosing a good physical plan

# Wrap up

- ▶ Cost estimation, towards choosing a good physical plan
- ▶ After the break
  - ▶ Answering queries using views



# Views

# Where we've been

cost estimation using data statistics

# Where we're headed

query optimization:

- ▶ logical optimization
- ▶ physical optimization

# Where we're headed

query optimization:

- ▶ logical optimization
- ▶ physical optimization
- ▶ next: the creation, maintenance, and use of “views”

# Views

Virtual/derived relations, providing alternative logical schemata


# Views

Virtual/derived relations, providing alternative logical schemata

Often desirable to provide

- ▶ security
- ▶ efficiency
- ▶ logical data independence

# Data independence



# Views

We've already studied a few special cases:

- ▶ Histograms/statistics as views




# Views


We've already studied a few special cases:

- ▶ Histograms/statistics as views
- ▶ Indexes as (family of) views

# Indexes as views



# Indexes as views



$\text{predicate}(r)$  vs.  $\text{predicate}(r')$

# Views

We've already studied a few special cases:

- ▶ Histograms/statistics as views
- ▶ Indexes as (family of) views
  - ▶ trees as hierarchical histograms: annotate nodes with actual or approximate counts of items in subtrees

# Views

We've already studied a few special cases:

- ▶ Histograms/statistics as views
- ▶ Indexes as (family of) views
  - ▶ trees as hierarchical histograms: annotate nodes with actual or approximate counts of items in subtrees
- ▶ GMAPs

# Views

In SQL, views are created with the CREATE VIEW statement

- ▶ `CREATE VIEW view_name AS expression`

# Views

In SQL, views are created with the CREATE VIEW statement

- ▶ CREATE VIEW *view\_name* AS *expression*

- ▶ for example

```
CREATE VIEW mng_view AS  
SELECT name, address, phone  
FROM emp  
WHERE title='manager'
```

# Views

Basic issues with views:



# Views

Basic issues with views:

- ▶ creation and implementation
- ▶ maintenance under updates
- ▶ answering queries using views: query containment


# Updating views

# Views

$I_B$

“Base” instance  $I_B$  and view  $v$

# Views



“Base” instance  $I_B$  and view  $v$


# Views

- ▶ We'd like to allow users to treat the view instance  $I_V$  just like any other (base) relation.

# Views


- ▶ We'd like to allow users to treat the view instance  $I_V$  just like any other (base) relation.
- ▶ in particular, it would be nice to support both queries and updates on  $I_V$

# Updating a view



Given update  $\delta_V$  on view instance  $I_V$ , find appropriate update  $\delta_B$  on the base data  $I_B$

# Updating a view



Given update  $\delta_V$  on view instance  $I_V$ , find appropriate update  $\delta_B$  on the base data  $I_B$



# Updating a view

Unfortunately, not all views are directly updateable

- ▶ `INSERT INTO mng_view  
VALUES (Fred, Eindhoven, 1234567)`

# Updating a view

Unfortunately, not all views are directly updateable

- ▶ `INSERT INTO mng_view`  
`VALUES (Fred, Eindhoven, 1234567)`
- ▶ how to update the EMP table, if it has a salary field? What should Fred's salary be?

# Updating a view

Another example

Suppose we have base table  $edge(x, y)$ , defining a directed graph, and the view

$$hop(x, y) = edge \bowtie edge$$

(i.e., paths of length 2)

# Updating a view

Another example

Suppose we have base table  $edge(x, y)$ , defining a directed graph, and the view

$$hop(x, y) = edge \bowtie edge$$

(i.e., paths of length 2)

How should we handle an insert on  $hop$ ? How should we materialize this in  $edge$ ?

# Updating a view

this is an active research topic!

- ▶ Can couple all view definitions with an appropriate “update policy”, using some formalism (e.g., so-called lenses)

# Updating a view

this is an active research topic!

- ▶ Can couple all view definitions with an appropriate “update policy”, using some formalism (e.g., so-called lenses)
- ▶ DBA can build trigger on view, to enforce “reasonable” behavior

# Updating a view

this is an active research topic!

- ▶ Can couple all view definitions with an appropriate “update policy”, using some formalism (e.g., so-called lenses)
- ▶ DBA can build trigger on view, to enforce “reasonable” behavior

```
CREATE TRIGGER mng_trigger
INSTEAD OF INSERT ON mng_view
BEGIN
  INSERT INTO emp VALUES
  (NEW.name, NEW.address, NEW.phone,
   'manager', $0);
END;
```

# Updating a view

Can proceed with simple restrictions



# Updating a view

Can proceed with simple restrictions

A view is *updateable* if

- ▶ FROM clause has only one relation,

# Updating a view

Can proceed with simple restrictions

A view is *updateable* if

- ▶ FROM clause has only one relation,
- ▶ SELECT clause contains only attributes (no expressions, etc.),

# Updating a view

Can proceed with simple restrictions

A view is *updateable* if

- ▶ FROM clause has only one relation,
- ▶ SELECT clause contains only attributes (no expressions, etc.),
- ▶ any attribute not listed in SELECT can be set to NULL, and

# Updating a view

Can proceed with simple restrictions

A view is *updateable* if

- ▶ FROM clause has only one relation,
- ▶ SELECT clause contains only attributes (no expressions, etc.),
- ▶ any attribute not listed in SELECT can be set to NULL, and
- ▶ no GROUP BY or HAVING clauses.


# Implementing views

# Implementing views

Two alternatives for view implementation


- ▶ virtual views: unfold any use of views in a query

# Virtual view unfolding



expression parse tree, using views  $V_1$  and  $V_2$

# Virtual view unfolding



expression parse tree, with  $V_1$  and  $V_2$  replaced with their definitions



# Virtual view unfolding

For example

$$\sigma_{address=Eindhoven}(MngView)$$

# Virtual view unfolding

For example

$$\sigma_{address=Eindhoven}(MngView)$$

becomes

$$\sigma_{address=Eindhoven}(\pi_{address,name,phone}(\sigma_{title=Manager}(Emp)))$$

# Views

Two alternatives for implementation

- ▶ virtual views: unfold any use of views in a query

# Views

Two alternatives for implementation

- ▶ virtual views: unfold any use of views in a query
- ▶ materialization

# Materialized views

- ▶ precompute and store view results

# Materialized views

- ▶ precompute and store view results
- ▶ more efficient for frequently used and/or expensive views
  - ▶ ex: the ATM view of your bank account

# Materialized views

- ▶ precompute and store view results
- ▶ more efficient for frequently used and/or expensive views
  - ▶ ex: the ATM view of your bank account
  - ▶ ex: view which denormalizes (i.e., incurs costly joins)

# Materialized views

Apps

- ▶ query optimization



# Materialized views

## Apps

- ▶ query optimization
- ▶ data warehousing
  - ▶ integration of data
  - ▶ OLAP

# Materialized views

## Apps

- ▶ query optimization
- ▶ data warehousing
  - ▶ integration of data
  - ▶ OLAP
- ▶ data replication/archiving

# Materialized views

## Apps

- ▶ query optimization
- ▶ data warehousing
  - ▶ integration of data
  - ▶ OLAP
- ▶ data replication/archiving
- ▶ data visualization

# Materialized views

## Apps

- ▶ query optimization
- ▶ data warehousing
  - ▶ integration of data
  - ▶ OLAP
- ▶ data replication/archiving
- ▶ data visualization
- ▶ caching in networked devices

# Materialized views

## Apps


- ▶ query optimization
- ▶ data warehousing
  - ▶ integration of data
  - ▶ OLAP
- ▶ data replication/archiving
- ▶ data visualization
- ▶ caching in networked devices
- ▶ ....

# Maintaining materialized views

# Materialized views

*Issue:* we must keep the view up-to-date, as base data evolves ...


# View maintenance



Given update  $\delta_B$  on base instance  $I_B$ , find appropriate update  $\delta_V$  on materialized view  $I_V$



# View maintenance



Given update  $\delta_B$  on base instance  $I_B$ , find appropriate update  $\delta_V$  on materialized view  $I_V$

# Materialized views: maintenance

- ▶ incremental vs. complete
- ▶ immediate vs. deferred

# Materialized views: maintenance

- ▶ manual code

# Materialized views: maintenance

- ▶ manual code
- ▶ triggers on base relations

# Materialized views: maintenance

- ▶ manual code
- ▶ triggers on base relations

```
CREATE TRIGGER mng_update
INSERT ON emp
BEGIN
    INSERT INTO mng_view VALUES
        (NEW.name, NEW.address, NEW.phone);
END;
```

# Materialized views: maintenance

The Counting Algorithm for incremental maintenance

# Materialized views: maintenance

The Counting Algorithm for incremental maintenance

- ▶ keep track of the number of derivations of a tuple in the view
  - ▶ essentially, this is the number of duplicates of the tuple in bag-evaluation of the view query

# Materialized views: maintenance

The Counting Algorithm for incremental maintenance

- ▶ keep track of the number of derivations of a tuple in the view
  - ▶ essentially, this is the number of duplicates of the tuple in bag-evaluation of the view query
- ▶ calculate update differential  $\delta_V$  for the view
  - ▶ e.g., for view  $V = R \bowtie S$ , and update  $\delta_R^+$  on  $R$ , we have  $\delta_V = \delta_R^+ \bowtie S$ , and

$$V' = V \cup \delta_V$$



# Materialized views: maintenance

The Counting Algorithm for incremental maintenance

- ▶ keep track of the number of derivations of a tuple in the view
  - ▶ essentially, this is the number of duplicates of the tuple in bag-evaluation of the view query
- ▶ calculate update differential  $\delta_V$  for the view
  - ▶ e.g., for view  $V = R \bowtie S$ , and update  $\delta_R^+$  on  $R$ , we have  $\delta_V = \delta_R^+ \bowtie S$ , and

$$V' = V \cup \delta_V$$

- ▶ rules for other algebra operators given in our textbook

# Materialized views: maintenance

## The Counting Algorithm

- ▶ example: for  
 $edge = \{(a, b), (b, c), (b, e), (a, d), (d, c)\}$ , we  
have  $hop = \{(a, c), (a, e)\}$

# Materialized views: maintenance

## The Counting Algorithm

- ▶ example: for  
 $edge = \{(a, b), (b, c), (b, e), (a, d), (d, c)\}$ , we  
have  $hop = \{(a, c), (a, e)\}$ 
  - ▶ how many ways is  $hop(a, c)$  derivable?

# Materialized views: maintenance

## The Counting Algorithm

- ▶ example: for  
 $edge = \{(a, b), (b, c), (b, e), (a, d), (d, c)\}$ , we  
have  $hop = \{(a, c), (a, e)\}$ 
  - ▶ how many ways is  $hop(a, c)$  derivable?
  - ▶ how many ways is  $hop(a, e)$  derivable?

# Materialized views: maintenance

## The Counting Algorithm

- ▶ example: for  
 $edge = \{(a, b), (b, c), (b, e), (a, d), (d, c)\}$ , we  
have  $hop = \{(a, c), (a, e)\}$ 
  - ▶ how many ways is  $hop(a, c)$  derivable?
  - ▶ how many ways is  $hop(a, e)$  derivable?
  - ▶ now, suppose  $\delta_{edge}^- = \{(a, b)\}$ .
  - ▶ then,  $\delta_{hop}^- = \{(a, c), (a, e)\}$ .

# Materialized views: maintenance

## The Counting Algorithm

- ▶ example: for  
 $edge = \{(a, b), (b, c), (b, e), (a, d), (d, c)\}$ , we have  $hop = \{(a, c), (a, e)\}$ 
  - ▶ how many ways is  $hop(a, c)$  derivable?
  - ▶ how many ways is  $hop(a, e)$  derivable?
  - ▶ now, suppose  $\delta_{edge}^- = \{(a, b)\}$ .
  - ▶ then,  $\delta_{hop}^- = \{(a, c), (a, e)\}$ .
- ▶ we know that after the update,  $hop = \{(a, c)\}$ .
- ▶ how to incrementally apply  $\delta_{hop}$  to get the correct  $hop$ ?

# Materialized views: maintenance

## The Counting Algorithm

- ▶ example: for  
 $edge = \{(a, b), (b, c), (b, e), (a, d), (d, c)\}$ , we have  $hop = \{(a, c), (a, e)\}$ 
  - ▶ how many ways is  $hop(a, c)$  derivable?
  - ▶ how many ways is  $hop(a, e)$  derivable?
  - ▶ now, suppose  $\delta_{edge}^- = \{(a, b)\}$ .
  - ▶ then,  $\delta_{hop}^- = \{(a, c), (a, e)\}$ .
- ▶ we know that after the update,  $hop = \{(a, c)\}$ .
- ▶ how to incrementally apply  $\delta_{hop}$  to get the correct  $hop$ ?
- ▶ check the counts of the elements of the view!

# Materialized views: maintenance

## The Counting Algorithm

- ▶ keep track of the number of derivations of a tuple in the view
- ▶ calculate update differential  $\delta_V$  for the view



# Materialized views: maintenance

## The Counting Algorithm

- ▶ keep track of the number of derivations of a tuple in the view
- ▶ calculate update differential  $\delta_V$  for the view
- ▶ for insertions (i.e.,  $\delta^+$ ), increment counts; for deletions (i.e.,  $\delta^-$ ), decrement counts
- ▶ if the count of a tuple goes to zero, remove it from the view

# Materialized views: maintenance

## The Counting Algorithm

- ▶ keep track of the number of derivations of a tuple in the view
- ▶ calculate update differential  $\delta_V$  for the view
- ▶ for insertions (i.e.,  $\delta^+$ ), increment counts; for deletions (i.e.,  $\delta^-$ ), decrement counts
- ▶ if the count of a tuple goes to zero, remove it from the view

in our example  $hop = \{(a, c) : 2, (a, e) : 1\}$ , and  
 $\delta_{hop}(hop) = \{(a, c) : 1\}$

# Using views

# Query containment & equivalence

When is a view useful for a given query?

# Query containment & equivalence

When is a view useful for a given query?

- ▶ more generally, given two queries  $Q_1$  and  $Q_2$ , is it the case that  $Q_1(I) \subseteq Q_2(I)$ , for any database instance  $I$ ?

# Query containment & equivalence

When is a view useful for a given query?

- ▶ more generally, given two queries  $Q_1$  and  $Q_2$ , is it the case that  $Q_1(I) \subseteq Q_2(I)$ , for any database instance  $I$ ?
- ▶ in this case we say  $Q_1$  is **contained** in  $Q_2$ , denoted  $Q_1 \subseteq Q_2$

# Query containment & equivalence

When is a view useful for a given query?

- ▶ more generally, given two queries  $Q_1$  and  $Q_2$ , is it the case that  $Q_1(I) \subseteq Q_2(I)$ , for any database instance  $I$ ?
- ▶ in this case we say  $Q_1$  is **contained** in  $Q_2$ , denoted  $Q_1 \subseteq Q_2$
- ▶ if  $Q_1 \subseteq Q_2$  and  $Q_2 \subseteq Q_1$ , then we say  $Q_1$  and  $Q_2$  are **equivalent**

# Query containment & equivalence

- recall the *hop* view, which we now write as

$$\text{hop}(x, y) \leftarrow \text{edge}(x, z), \text{edge}(z, y)$$

in familiar datalog notation



# Query containment & equivalence

- ▶ recall the *hop* view, which we now write as

$$\text{hop}(x, y) \leftarrow \text{edge}(x, z), \text{edge}(z, y)$$

in familiar datalog notation

- ▶ next, consider

$$\text{hop}'(v, w) \leftarrow \text{edge}(v, u), \text{edge}(u, w), \text{edge}(w, w)$$

# Query containment & equivalence

- ▶ recall the *hop* view, which we now write as

$$\text{hop}(x, y) \leftarrow \text{edge}(x, z), \text{edge}(z, y)$$

in familiar datalog notation

- ▶ next, consider

$$\text{hop}'(v, w) \leftarrow \text{edge}(v, u), \text{edge}(u, w), \text{edge}(w, w)$$

- ▶ is it the case that  $\text{hop} \subseteq \text{hop}'$ ?
- ▶ is it the case that  $\text{hop}' \subseteq \text{hop}$ ?

# Query containment & equivalence

- ▶ recall the *hop* view, which we now write as

$$\text{hop}(x, y) \leftarrow \text{edge}(x, z), \text{edge}(z, y)$$

in familiar datalog notation

- ▶ next, consider

$$\text{hop}'(v, w) \leftarrow \text{edge}(v, u), \text{edge}(u, w), \text{edge}(w, w)$$

- ▶ is it the case that  $\text{hop} \subseteq \text{hop}'$ ?
- ▶ is it the case that  $\text{hop}' \subseteq \text{hop}$ ?
- ▶ how can we prove this?
  - ▶ we will restrict our discussion now to conjunctive queries (containment is undecidable for FO ...)

# Query containment & equivalence

## Homomorphisms

- ▶ a mapping from the variables of  $Q_2$  to the variables of  $Q_1$ , such that
  - ▶ the head of  $Q_2$  becomes the head of  $Q_1$
  - ▶ each subgoal of  $Q_2$  becomes some subgoal of  $Q_1$

# Query containment & equivalence

## Homomorphisms

- ▶ a mapping from the variables of  $Q_2$  to the variables of  $Q_1$ , such that
  - ▶ the head of  $Q_2$  becomes the head of  $Q_1$
  - ▶ each subgoal of  $Q_2$  becomes some subgoal of  $Q_1$
- ▶ it isn't necessary for every subgoal of  $Q_1$  to be mapped onto

# Query containment & equivalence

for example, the homomorphism  $\varphi$  defined as

$$x \rightarrow v,$$

$$y \rightarrow w,$$

$$z \rightarrow u$$

# Query containment & equivalence

for example, the homomorphism  $\varphi$  defined as

$$x \rightarrow v,$$

$$y \rightarrow w,$$

$$z \rightarrow u$$

maps

$$\text{hop}(x, y) \leftarrow \text{edge}(x, z), \text{edge}(z, y)$$

onto

$$\text{hop}'(v, w) \leftarrow \text{edge}(v, u), \text{edge}(u, w), \text{edge}(w, w)$$

# Query containment & equivalence

## Theorem

$Q_1 \subseteq Q_2$  if and only if there exists a homomorphism from  $Q_2$  to  $Q_1$



# Query containment & equivalence

## Theorem

$Q_1 \subseteq Q_2$  if and only if there exists a homomorphism from  $Q_2$  to  $Q_1$

**proof: (if)** Let  $\varphi : Q_2 \rightarrow Q_1$  be a homomorphism, and let  $I$  be a database. Every tuple  $t \in Q_1(I)$  is produced by some substitution  $\psi_t$  on the variables of  $Q_1$  that make all of  $Q_1$ 's subgoals true in  $I$ . Then  $\psi_t \circ \varphi$  is a substitution for variables of  $Q_2$  such that  $t \in Q_2(I)$ . Hence,  $Q_1 \subseteq Q_2$ .

# Query containment & equivalence

## Theorem

$Q_1 \subseteq Q_2$  if and only if there exists a homomorphism from  $Q_2$  to  $Q_1$

**proof, continued: (only if)** Create a fresh unique atom for each variable of  $Q_1$ , and let  $I_{Q_1}$  be the database instance consisting of all the subgoals of  $Q_1$ , with the chosen atoms substituted for variables. Now, note that  $Q_1(I_{Q_1})$  contains the “atom-head”  $t$  of  $Q_1$ . Since  $Q_1 \subseteq Q_2$ , it must also be that  $t \in Q_2(I_{Q_1})$ .

# Query containment & equivalence

## Theorem

$Q_1 \subseteq Q_2$  if and only if there exists a homomorphism from  $Q_2$  to  $Q_1$

**proof, continued:** Let  $\psi_t$  be the substitution of constants from  $I_{Q_1}$  for the variables of  $Q_2$  that makes each subgoal of  $Q_2$  a tuple of the instance  $I_{Q_1}$  and yields  $t$  as the head; and, let  $\varphi$  be the substitution that maps constants of  $I_{Q_1}$  to their unique corresponding variable of  $Q_1$ . Then  $\varphi \circ \psi_t$  is a homomorphism from  $Q_2$  to  $Q_1$ . □

# Query containment & equivalence

unfortunately, checking containment for conjunctive queries is an NP-complete problem

# Query containment & equivalence

unfortunately, checking containment for conjunctive queries is an NP-complete problem

1. given a mapping  $m$  from  $Q_2$  to  $Q_1$ , we can check if  $m$  is indeed a homomorphism in polynomial time

# Query containment & equivalence

unfortunately, checking containment for conjunctive queries is an NP-complete problem

1. given a mapping  $m$  from  $Q_2$  to  $Q_1$ , we can check if  $m$  is indeed a homomorphism in polynomial time
2. Let  $G = (V, E)$  be a graph and  $k$  be an integer. Consider, for set  $C$  of  $k$  new distinct variables,

$$Q_2 = out() \leftarrow \bigwedge_{(u,v) \in E} E(u, v)$$

$$Q_1 = out() \leftarrow \bigwedge_{u,v \in C, u \neq v} E(u, v)$$

Then  $Q_1 \subseteq Q_2$  iff  $G$  has a  $k$ -coloring.



# Query containment & equivalence

Fortunately, queries are often quite small, especially with respect to the size of data

Furthermore, checking query containment for **acyclic** conjunctive queries is *tractable* (i.e., computable in polynomial time). More on this in a later lecture ...

## Exercise: answering queries with views

Consider the following conjunctive query.

$$Q : \text{result}(A) \leftarrow r(A, B), r(A, C), s(B, D, E), s(B, F, F)$$

Minimize  $Q$ . In other words, give a query  $Q'$  that (i) has the smallest possible body and (ii) is equivalent to  $Q$ . Demonstrate that your query satisfies both of these properties.



Wrap Up

# Wrap up

- ▶ Cost estimation, towards choosing a good physical plan
- ▶ the construction, maintenance, and use of views

# Wrap up

- ▶ Cost estimation, towards choosing a good physical plan
- ▶ the construction, maintenance, and use of views
- ▶ Next time
  - ▶ putting everything together for query optimization

# Wrap up

- ▶ Cost estimation, towards choosing a good physical plan
- ▶ the construction, maintenance, and use of views
- ▶ Next time
  - ▶ putting everything together for query optimization
- ▶ **Reminder:** team project report due by Wednesday 13 May

# Credits

Ullman 1999