

# Query optimization

## Lecture 7

2ID35, Spring 2015

George Fletcher

Faculteit Wiskunde & Informatica  
Technische Universiteit Eindhoven

13 May 2015

# Admin

- ▶ Project part 2 due today
  - ▶ Part 3 has been posted

# Admin

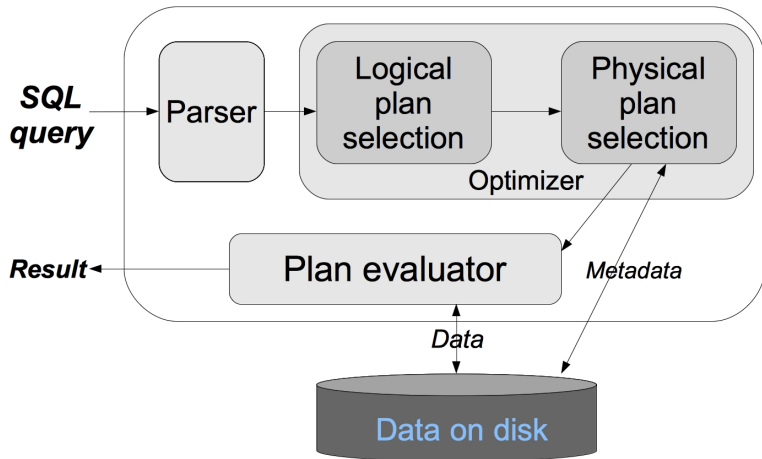
- ▶ Project part 2 due today
  - ▶ Part 3 has been posted
- ▶ No lecture this Friday or next Wednesday
- ▶ Written individual assignment will be posted this week

# Last time

Size estimation (heuristics, histograms)

View creation, maintenance, and use

# The life of a query

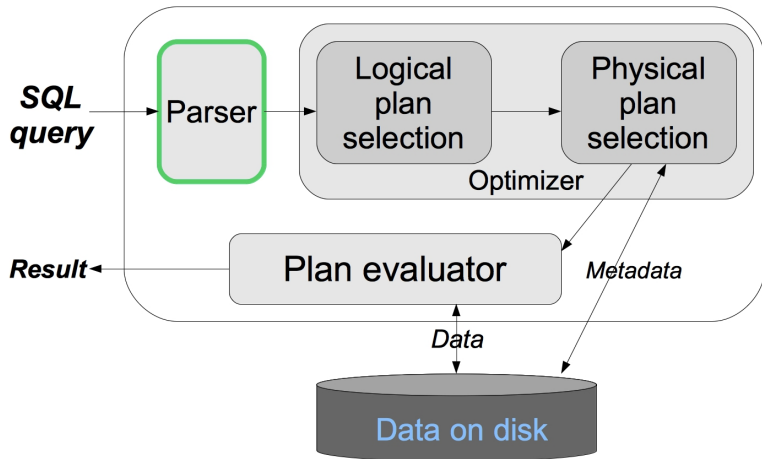


# Today

## Query compilation

- ▶ parsing
- ▶ logical optimization
- ▶ physical optimization

# The life of a query: parsing



# The life of a query: parsing

- ▶ Employee(EID, Name, ECity)
- ▶ Company(CID, Name, CCity)
- ▶ WorksFor(EID, CID, Salary)

*What are the ID's of employees living and working in Best with above-average salaries?*

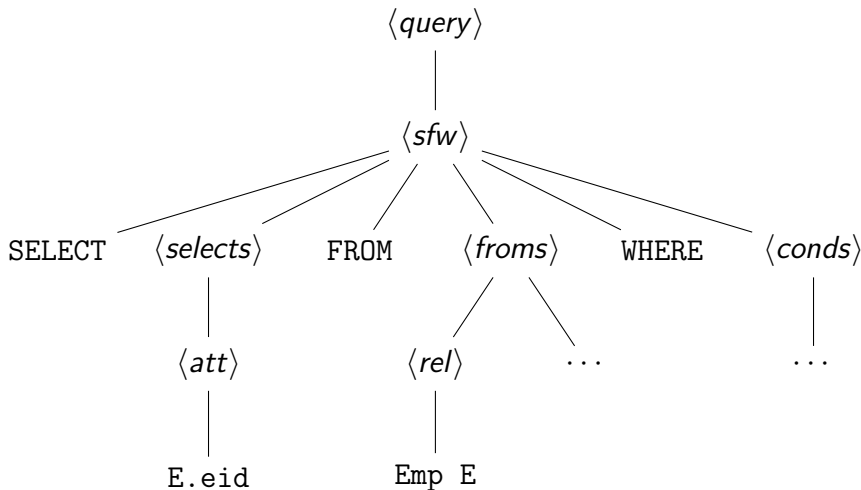


# The life of a query: parsing

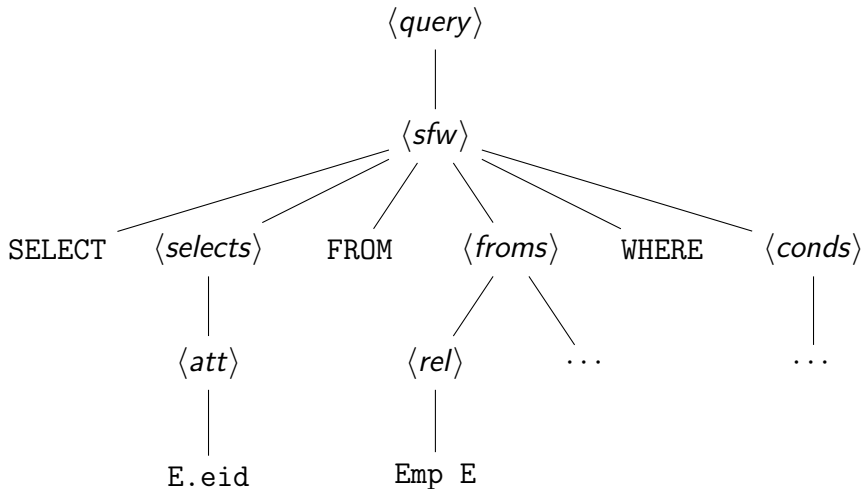
- ▶ Employee(EID, Name, ECity)
- ▶ Company(CID, Name, CCity)
- ▶ WorksFor(EID, CID, Salary)

```
SELECT E.EID
FROM Employee E, WorksFor W, Company C
WHERE E.EID = W.EID AND W.CID = C.CID
      AND E.ECity = 'Best'
      AND C.CCity = 'Best'
      AND W.Salary > (SELECT AVG(SALARY)
                      FROM WorksFor)
```

# The life of a query: parsing



# The life of a query: parsing



... and preprocess for semantics (e.g., typing)

# The life of a query: parsing

Parser generates a collection of query *blocks*

- ▶ block = a S-F-W query with no nesting
  - ▶ essentially, a conjunctive query

# The life of a query: parsing

Parser generates a collection of query *blocks*

- ▶ block = a S-F-W query with no nesting
  - ▶ essentially, a conjunctive query
- ▶ typically focus on optimizing one block at a time

# The life of a query: parsing

- ▶ Employee(EID, Name, ECity)
- ▶ Company(CID, Name, CCity)
- ▶ WorksFor(EID, CID, Salary)

```
SELECT E.EID
FROM Employee E, WorksFor W, Company C
WHERE E.EID = W.EID AND W.CID = C.CID
      AND E.ECity = 'Best'
      AND C.CCity = 'Best'
      AND W.Salary >
      (SELECT AVG(SALARY)
       FROM WorksFor)
```

nested block

# The life of a query: parsing

- ▶ Employee(EID, Name, ECity)
- ▶ Company(CID, Name, CCity)
- ▶ WorksFor(EID, CID, Salary)

```
SELECT E.EID
FROM Employee E, WorksFor W, Company C
WHERE E.EID = W.EID AND W.CID = C.CID
      AND E.ECity = 'Best'
      AND C.CCity = 'Best'
      AND W.Salary >
              (SELECT AVG(SALARY)
               FROM WorksFor)
```

outer block

# The life of a query: parsing

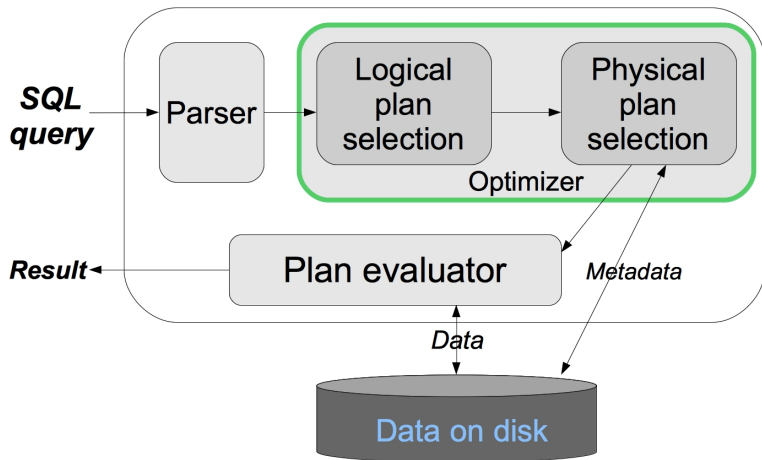
- ▶ Employee(EID, Name, ECity)
- ▶ Company(CID, Name, CCity)
- ▶ WorksFor(EID, CID, Salary)

```
SELECT E.EID
FROM Employee E, WorksFor W, Company C
WHERE E.EID = W.EID AND W.CID = C.CID
      AND E.ECity = 'Best'
      AND C.CCity = 'Best'
      AND W.Salary >
      (SELECT AVG(SALARY)
       FROM WorksFor)
```

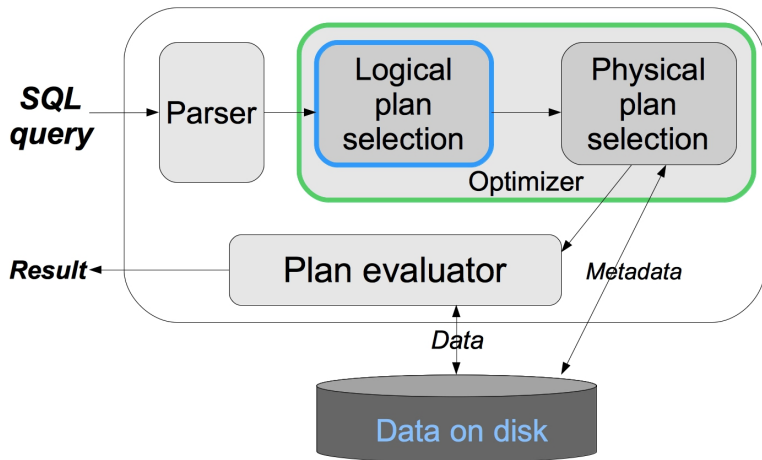
focus on optimizing each block separately



# The life of a query: optimization



# The life of a query: logical optimization



# The life of a query: logical optimization

*Goal:* map a query-block to a “preferred” logical query plan

# The life of a query: logical optimization

*Goal:* map a query-block to a “preferred” logical query plan

*Step A:* map the block's parse tree to equivalent RA expression, working bottom-up

1. map FROM list to cartesian product of all relations

# The life of a query: logical optimization

*Goal:* map a query-block to a “preferred” logical query plan

*Step A:* map the block’s parse tree to equivalent RA expression, working bottom-up

1. map FROM list to cartesian product of all relations
2. form single  $\sigma$  for WHERE conditions

# The life of a query: logical optimization

*Goal:* map a query-block to a “preferred” logical query plan

*Step A:* map the block’s parse tree to equivalent RA expression, working bottom-up

1. map FROM list to cartesian product of all relations
2. form single  $\sigma$  for WHERE conditions
3. form  $\pi$  list from the SELECT clause

# The life of a query: logical optimization


*Goal:* map a query-block to a “preferred” logical query plan

*Step A:* map the block's parse tree to equivalent RA expression, working bottom-up

1. map FROM list to cartesian product of all relations
2. form single  $\sigma$  for WHERE conditions
3. form  $\pi$  list from the SELECT clause


any aggregation can be applied after this  $\pi \cdot \sigma \cdot \times$

# The life of a query: logical optimization






# The life of a query: logical optimization



# The life of a query: logical optimization



- (1)  $e.eid = w.eid$
- (2)  $w.cid = c.cid$
- (3)  $e.ecity = \text{'Best'}$
- (4)  $c.ccity = \text{'Best'}$
- (5)  $w.salary > \mathcal{V}$

# The life of a query: logical optimization

*Step B:* improve this initial logical query plan (i.e., generate a preferred logical query plan)

# The life of a query: logical optimization

*Step B:* improve this initial logical query plan (i.e., generate a preferred logical query plan)

based on equivalence rules, used to identify different ways of formulating the query

# The life of a query: logical optimization

*Step B:* improve this initial logical query plan (i.e., generate a preferred logical query plan)

based on equivalence rules, used to identify different ways of formulating the query

this defines a search space of alternative plans, to be considered by the optimizer

# RA equivalence rules: commutativity & associativity

$$\begin{aligned}R \star S &= S \star R \\ (R \star S) \star T &= R \star (S \star T)\end{aligned}$$

for  $\star \in \{\times, \bowtie, \cup, \cap\}$

# RA equivalence rules: commutativity & associativity

So, we have for example

$$\begin{aligned}R \cap S &= S \cap R \\ (R \cap S) \cap T &= R \cap (S \cap T)\end{aligned}$$

# RA equivalence rules: commutativity & associativity

So, we have for example

$$\begin{aligned}R \cap S &= S \cap R \\ (R \cap S) \cap T &= R \cap (S \cap T)\end{aligned}$$

*example application.*

$$\begin{aligned}\sigma_{\theta}(\pi_A(R)) \cap (\pi_A(S) \cap \pi_A(T)) \\ = (\sigma_{\theta}(\pi_A(R)) \cap \pi_A(S)) \cap \pi_A(T)\end{aligned}$$



# RA equivalence rules: selection laws

splitting

$$\sigma_{C_1 \wedge \dots \wedge C_n}(R) = \sigma_{C_1}(\dots \sigma_{C_n}(R) \dots)$$

# RA equivalence rules: selection laws

splitting

$$\sigma_{C_1 \wedge \dots \wedge C_n}(R) = \sigma_{C_1}(\dots \sigma_{C_n}(R) \dots)$$

$$\sigma_{C_1 \vee C_2}(R) = \sigma_{C_1}(R) \cup \sigma_{C_2}(R)$$

# RA equivalence rules: selection laws

splitting

$$\sigma_{C_1 \wedge \dots \wedge C_n}(R) = \sigma_{C_1}(\dots \sigma_{C_n}(R) \dots)$$

$$\sigma_{C_1 \vee C_2}(R) = \sigma_{C_1}(R) \cup \sigma_{C_2}(R)$$

*example application.*

$$\sigma_{a=1}(\sigma_{b<c}(R)) \cup \sigma_{a=3}(\sigma_{b<c}(R))$$

# RA equivalence rules: selection laws

splitting

$$\sigma_{C_1 \wedge \dots \wedge C_n}(R) = \sigma_{C_1}(\dots \sigma_{C_n}(R) \dots)$$

$$\sigma_{C_1 \vee C_2}(R) = \sigma_{C_1}(R) \cup \sigma_{C_2}(R)$$

*example application.*

$$\sigma_{a=1}(\sigma_{b<c}(R)) \cup \sigma_{a=3}(\sigma_{b<c}(R)) = \sigma_{a=1 \vee a=3}(\sigma_{b<c}(R))$$

# RA equivalence rules: selection laws

splitting

$$\sigma_{C_1 \wedge \dots \wedge C_n}(R) = \sigma_{C_1}(\dots \sigma_{C_n}(R) \dots)$$

$$\sigma_{C_1 \vee C_2}(R) = \sigma_{C_1}(R) \cup \sigma_{C_2}(R)$$

*example application.*

$$\begin{aligned} \sigma_{a=1}(\sigma_{b<c}(R)) \cup \sigma_{a=3}(\sigma_{b<c}(R)) &= \sigma_{a=1 \vee a=3}(\sigma_{b<c}(R)) \\ &= \sigma_{(a=1 \vee a=3) \wedge b<c}(R) \end{aligned}$$

# RA equivalence rules: selection laws

union

$$\sigma_{\theta}(R \cup S) = \sigma_{\theta}(R) \cup \sigma_{\theta}(S)$$

# RA equivalence rules: selection laws

union

$$\sigma_{\theta}(R \cup S) = \sigma_{\theta}(R) \cup \sigma_{\theta}(S)$$

difference

$$\sigma_{\theta}(R - S) = \sigma_{\theta}(R) - \sigma_{\theta}(S)$$

# RA equivalence rules: selection laws

union

$$\sigma_{\theta}(R \cup S) = \sigma_{\theta}(R) \cup \sigma_{\theta}(S)$$

difference

$$\begin{aligned}\sigma_{\theta}(R - S) &= \sigma_{\theta}(R) - \sigma_{\theta}(S) \\ &= \sigma_{\theta}(R) - S\end{aligned}$$



# RA equivalence rules: selection laws

union

$$\sigma_{\theta}(R \cup S) = \sigma_{\theta}(R) \cup \sigma_{\theta}(S)$$

difference

$$\begin{aligned}\sigma_{\theta}(R - S) &= \sigma_{\theta}(R) - \sigma_{\theta}(S) \\ &= \sigma_{\theta}(R) - S\end{aligned}$$

and similarly for  $\times, \bowtie, \cap$

# RA equivalence rules: selection laws

union

$$\sigma_{\theta}(R \cup S) = \sigma_{\theta}(R) \cup \sigma_{\theta}(S)$$

difference

$$\begin{aligned}\sigma_{\theta}(R - S) &= \sigma_{\theta}(R) - \sigma_{\theta}(S) \\ &= \sigma_{\theta}(R) - S\end{aligned}$$

and similarly for  $\times, \bowtie, \cap$

proof of the difference rule

# RA equivalence rules: selection laws

*example application.* for  $R(a, b)$  and  $S(b, c)$

$$\sigma_{(a=1 \vee a=3) \wedge b < c}(R \bowtie S)$$

# RA equivalence rules: selection laws

*example application.* for  $R(a, b)$  and  $S(b, c)$

$$\sigma_{(a=1 \vee a=3) \wedge b < c}(R \bowtie S) = \sigma_{a=1 \vee a=3}(\sigma_{b < c}(R \bowtie S))$$

# RA equivalence rules: selection laws

*example application.* for  $R(a, b)$  and  $S(b, c)$

$$\begin{aligned}\sigma_{(a=1 \vee a=3) \wedge b < c}(R \bowtie S) &= \sigma_{a=1 \vee a=3}(\sigma_{b < c}(R \bowtie S)) \\ &= \sigma_{a=1 \vee a=3}(R \bowtie \sigma_{b < c}(S))\end{aligned}$$

# RA equivalence rules: selection laws

*example application.* for  $R(a, b)$  and  $S(b, c)$

$$\begin{aligned}\sigma_{(a=1 \vee a=3) \wedge b < c}(R \bowtie S) &= \sigma_{a=1 \vee a=3}(\sigma_{b < c}(R \bowtie S)) \\ &= \sigma_{a=1 \vee a=3}(R \bowtie \sigma_{b < c}(S)) \\ &= \sigma_{a=1 \vee a=3}(R) \bowtie \sigma_{b < c}(S).\end{aligned}$$

# RA equivalence rules: selection laws

*example application.* for  $R(a, b)$  and  $S(b, c)$

$$\begin{aligned}\sigma_{(a=1 \vee a=3) \wedge b < c}(R \bowtie S) &= \sigma_{a=1 \vee a=3}(\sigma_{b < c}(R \bowtie S)) \\ &= \sigma_{a=1 \vee a=3}(R \bowtie \sigma_{b < c}(S)) \\ &= \sigma_{a=1 \vee a=3}(R) \bowtie \sigma_{b < c}(S).\end{aligned}$$

note how this brings the selection conditions closer to the input relations ...

## RA equivalence rules: projection law

if, for each  $1 \leq i \leq n$ , we have  $a_i \subseteq a_{i+1}$ , for subsets  $a_1, \dots, a_n$  of  $atts(R)$ , then

$$\pi_{a_1}(R) = \pi_{a_1}(\pi_{a_2}(\cdots(\pi_{a_n}(R))\cdots)).$$



## RA equivalence rules: projection law

if, for each  $1 \leq i \leq n$ , we have  $a_i \subseteq a_{i+1}$ , for subsets  $a_1, \dots, a_n$  of  $atts(R)$ , then

$$\pi_{a_1}(R) = \pi_{a_1}(\pi_{a_2}(\dots(\pi_{a_n}(R))\dots)).$$

*example application.* for  $R(a, b, c)$

$$\pi_a(R) = \pi_a(\pi_{a,b}(R)).$$

# RA equivalence rules

*Prove or disprove:  $\pi_a(R - S) = \pi_a(R) - \pi_a(S)$ ,  
where  $a$  is a nonempty set of attributes in  $R$  and  $S$ .*

# RA equivalence rules

*Prove or disprove:*  $\pi_a(R - S) = \pi_a(R) - \pi_a(S)$ ,  
where  $a$  is a nonempty set of attributes in  $R$  and  $S$ .

*counterexample.* Consider  $R(a, b)$  and  $S(a, b)$ , with  
respective instances  $r = \{(1, 2)\}$  and  $s = \{(1, 3)\}$ .  
Then  $\pi_a(r - s) = \{(1)\}$ , but  $\pi_a(r) - \pi_a(s) = \{\}$ .

## RA equivalence rules

*Prove or disprove:*  $\pi_x(R \bowtie S) = \pi_x(\pi_{xy}(R) \bowtie S)$ ,  
where  $y = \text{atts}(R) \cap \text{atts}(S)$  and  $x \subseteq \text{atts}(R)$ .

# RA equivalence rules

*Prove or disprove:*  $\pi_x(R \bowtie S) = \pi_x(\pi_{xy}(R) \bowtie S)$ ,  
where  $y = \text{atts}(R) \cap \text{atts}(S)$  and  $x \subseteq \text{atts}(R)$ .

*proof.* We show that  $\pi_x(R \bowtie S) \subseteq \pi_x(\pi_{xy}(R) \bowtie S)$ . The other direction is similar.

Suppose that  $t \in \pi_x(R \bowtie S)$ . Then we have that (1) there exists  $t' \in R \bowtie S$  such that  $t = t'[x]$  (i.e.,  $t'$  projected on  $x$ ).

It then follows that there exists  $r \in R$  and  $s \in S$  such that (2)  $t'[\text{atts}(R)] = r$ , (3)  $t'[\text{atts}(S)] = s$ , and (4)  $r[y] = s[y]$ .

From (2) we have that (5)  $t'[xy] \in \pi_{xy}(R)$ .

From (2), (4), and (5) we have that (6)  $t'[xy] \bowtie s \in \pi_{xy}(R) \bowtie S$ .

From (1) and (6), we have that  $t = t'[x] \in \pi_x(\pi_{xy}(R) \bowtie S)$ , as desired.

# RA equivalence rules: SPJ laws

1. if  $atts(\theta) \subseteq \{a_1, \dots, a_n\}$ , then

$$\pi_{\{a_1, \dots, a_n\}}(\sigma_{\theta}(R)) = \sigma_{\theta}(\pi_{\{a_1, \dots, a_n\}}(R))$$

# RA equivalence rules: SPJ laws

1. if  $atts(\theta) \subseteq \{a_1, \dots, a_n\}$ , then

$$\pi_{\{a_1, \dots, a_n\}}(\sigma_{\theta}(R)) = \sigma_{\theta}(\pi_{\{a_1, \dots, a_n\}}(R))$$

2.  $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$

# RA equivalence rules: SPJ laws

3. if  $a_1 \cup a_2 = a$ ,  $a_1 \subseteq \text{atts}(R)$ , and  $a_2 \subseteq \text{atts}(S)$ , then

$$\pi_a(R \times S) = \pi_{a_1}(R) \times \pi_{a_2}(S)$$



## RA equivalence rules: SPJ laws

3. if  $a_1 \cup a_2 = a$ ,  $a_1 \subseteq \text{atts}(R)$ , and  $a_2 \subseteq \text{atts}(S)$ , then

$$\pi_a(R \times S) = \pi_{a_1}(R) \times \pi_{a_2}(S)$$

4. if  $a_1 \cup a_2 = a$ ,  $a_1 \subseteq \text{atts}(R)$ ,  $a_2 \subseteq \text{atts}(S)$ , and  $\text{atts}(\theta) \subseteq a$ , then

$$\pi_a(R \bowtie_{\theta} S) = \pi_{a_1}(R) \bowtie_{\theta} \pi_{a_2}(S)$$

## RA equivalence rules: SPJ laws

3. if  $a_1 \cup a_2 = a$ ,  $a_1 \subseteq \text{atts}(R)$ , and  $a_2 \subseteq \text{atts}(S)$ , then

$$\pi_a(R \times S) = \pi_{a_1}(R) \times \pi_{a_2}(S)$$

4. if  $a_1 \cup a_2 = a$ ,  $a_1 \subseteq \text{atts}(R)$ ,  $a_2 \subseteq \text{atts}(S)$ , and  $\text{atts}(\theta) \subseteq a$ , then

$$\pi_a(R \bowtie_{\theta} S) = \pi_{a_1}(R) \bowtie_{\theta} \pi_{a_2}(S)$$

exercise. prove 4

# RA equivalence rules

Some heuristics

1. push down selections ( $\sigma$ ) as far as they can go, splitting conjunctions in conditions ( $\wedge$ )

# RA equivalence rules

Some heuristics

1. push down selections ( $\sigma$ ) as far as they can go, splitting conjunctions in conditions ( $\wedge$ )
  - ▶ single most important strategy

# RA equivalence rules

Some heuristics


1. push down selections ( $\sigma$ ) as far as they can go, splitting conjunctions in conditions ( $\wedge$ )
  - ▶ single most important strategy
2. push down projections ( $\pi$ )

# RA equivalence rules

## Some heuristics

1. push down selections ( $\sigma$ ) as far as they can go, splitting conjunctions in conditions ( $\wedge$ )
  - ▶ single most important strategy
2. push down projections ( $\pi$ )
3. if possible, turn  $\sigma + \times$  into  $\bowtie$ , which is generally much cheaper to evaluate than  $\sigma$  and  $\times$  evaluated separately

# RA equivalence rules

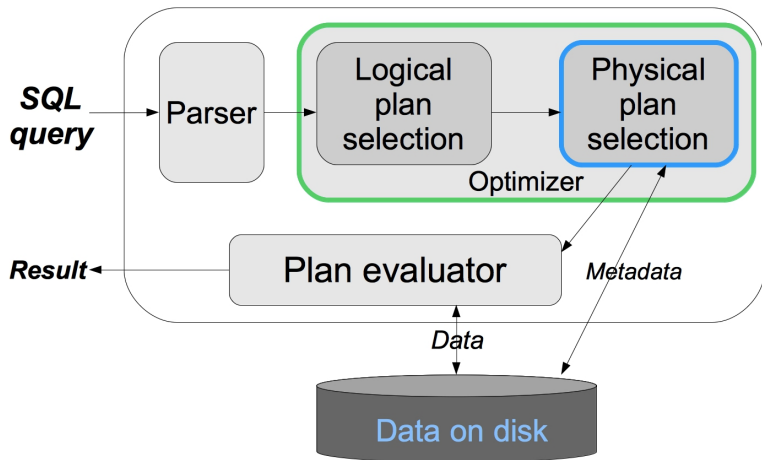


$e(eid, name, ecity)$   
 $c(cid, name, ccity)$   
 $w(eid, cid, salary)$

- (1)  $e.eid = w.eid$
- (2)  $w.cid = c.cid$
- (3)  $e.ecity = \text{'Best'}$
- (4)  $c.ccity = \text{'Best'}$
- (5)  $w.salary > \mathcal{V}$

let's logically optimize our query, under these heuristics

# The life of a query: physical optimization





# The life of a query: physical optimization

## Decisions

- ▶ two crucial decisions the optimizer makes
  - ▶ the order in which the physical operators are applied on the input relations
  - ▶ the choice of physical algorithms for each node of the plan
- ▶ of course, these are not independent

# The life of a query: physical optimization

## cost-based query optimization

- ▶ enumerate alternative plans, estimate the cost of each plan, pick the plan with minimum cost
- ▶ access methods and join algorithms define a search space
  - ▶ this space can be huge!
  - ▶ plan enumeration is the exploration of this search space

# The life of a query: physical optimization

For example, suppose we have five relations, only one access method, only one join algorithm, and we only consider “linear” joins

# The life of a query: physical optimization

For example, suppose we have five relations, only one access method, only one join algorithm, and we only consider “linear” joins

- ▶ there are then  $5! = 120$  possible plans

# The life of a query: physical optimization

For example, suppose we have five relations, only one access method, only one join algorithm, and we only consider “linear” joins

- ▶ there are then  $5! = 120$  possible plans
- ▶ if we add one additional access method, we then have  $2^5 \cdot 5! = 3840$  possible plans

# The life of a query: physical optimization

For example, suppose we have five relations, only one access method, only one join algorithm, and we only consider “linear” joins

- ▶ there are then  $5! = 120$  possible plans
- ▶ if we add one additional access method, we then have  $2^5 \cdot 5! = 3840$  possible plans
- ▶ if we add one additional join algorithm, we then have  $2^4 \cdot 2^5 \cdot 5! = 61440$  possible plans

# The life of a query: physical optimization

For example, suppose we have five relations, only one access method, only one join algorithm, and we only consider “linear” joins

- ▶ there are then  $5! = 120$  possible plans
- ▶ if we add one additional access method, we then have  $2^5 \cdot 5! = 3840$  possible plans
- ▶ if we add one additional join algorithm, we then have  $2^4 \cdot 2^5 \cdot 5! = 61440$  possible plans

In general, with  $n$  relations, there are  $\frac{(2(n-1))!}{(n-1)!}$  different join orders. For  $n = 5$ , this gives us 860,160 possible plans with two access methods and join algorithms. For  $n = 10$ , this gives us roughly 17.6 billion plans ...

# The life of a query: physical optimization

cost-based query optimization

- ▶ exhaustive search is certainly out of the question
- ▶ it could be that exploring the search space might take longer than actually evaluating the query



# The life of a query: physical optimization

## cost-based query optimization

- ▶ exhaustive search is certainly out of the question
- ▶ it could be that exploring the search space might take longer than actually evaluating the query
- ▶ the manner in which the plan space is explored describes a (physical) query optimization method
  - ▶ dynamic programming, rule-based optimization, randomized search, ...

# The life of a query: physical optimization

cost-based query optimization

- ▶ costing plans and exploring a plan space is nontrivial
- ▶ now, consider that the DBMS must do this for 1000's of queries, simultaneously!
- ▶ hence, all of this must be done quickly, without looking back

# The life of a query: physical optimization

## cost-based query optimization

- ▶ costing plans and exploring a plan space is nontrivial
- ▶ now, consider that the DBMS must do this for 1000's of queries, simultaneously!
- ▶ hence, all of this must be done quickly, without looking back
- ▶ query optimization is very much still an active area of research
- ▶ indeed, rarely will an optimizer find the optimal plan
  - ▶ it must, however, not pick a bad plan – just an OK one

# Query optimization: dynamic-programming algorithm

```
procedure FindBestPlan(S)  
  if (bestplan[S].cost  $\neq \infty$ ) /* bestplan[S] already computed */  
    return bestplan[S]  
  if (S contains only 1 relation)  
    set bestplan[S].plan and bestplan[S].cost based on best way of accessing S  
  else for each non-empty subset S1 of S such that S1  $\neq S$   
    P1 = FindBestPlan(S1)  
    P2 = FindBestPlan(S − S1)  
    A = best algorithm for joining results of P1 and P2  
    cost = P1.cost + P2.cost + cost of A  
    if cost < bestplan[S].cost  
      bestplan[S].cost = cost  
      bestplan[S].plan = “execute P1.plan; execute P2.plan;  
                           join results of P1 and P2 using A”  
return bestplan[S]
```

# Query optimization: dynamic-programming algorithm

```
procedure FindBestPlan(S)
  if (bestplan[S].cost  $\neq \infty$ ) /* bestplan[S] already computed */
    return bestplan[S]
  if (S contains only 1 relation)
    set bestplan[S].plan and bestplan[S].cost based on best way of accessing S
  else for each non-empty subset S1 of S such that S1  $\neq S$ 
    P1 = FindBestPlan(S1)
    P2 = FindBestPlan(S − S1)
    A = best algorithm for joining results of P1 and P2
    cost = P1.cost + P2.cost + cost of A
    if cost < bestplan[S].cost
      bestplan[S].cost = cost
      bestplan[S].plan = “execute P1.plan; execute P2.plan;  
                           join results of P1 and P2 using A”
  return bestplan[S]
```

Running time is  $\mathcal{O}(3^n)$ , which gives us roughly 59,000 plans for  $n = 10$  (compare this with 17.6 billion plans)

# Query optimization: dynamic-programming algorithm

```
procedure FindBestPlan(S)  
  if (bestplan[S].cost  $\neq \infty$ ) /* bestplan[S] already computed */  
    return bestplan[S]  
  if (S contains only 1 relation)  
    set bestplan[S].plan and bestplan[S].cost based on best way of accessing S  
  else for each non-empty subset S1 of S such that S1  $\neq S$   
    P1 = FindBestPlan(S1)  
    P2 = FindBestPlan(S - S1)  
    A = best algorithm for joining results of P1 and P2  
    cost = P1.cost + P2.cost + cost of A  
    if cost < bestplan[S].cost  
      bestplan[S].cost = cost  
      bestplan[S].plan = "execute P1.plan; execute P2.plan;  
                           join results of P1 and P2 using A"  
  return bestplan[S]
```

- ▶ most widely used strategy
- ▶ works well for queries with fewer than 10 to 15 joins

# Query optimization in System R

- ▶ follows the classic dynamic programming approach

# Query optimization in System R

- ▶ follows the classic dynamic programming approach
- ▶ heuristics: use the equivalence rules to push down selections and projections, and delay cartesian products



# Query optimization in System R

- ▶ follows the classic dynamic programming approach
- ▶ heuristics: use the equivalence rules to push down selections and projections, and delay cartesian products
- ▶ constraints: left-deep plans, nested-loops and sort-merge join only

# Query optimization in System R


- ▶ follows the classic dynamic programming approach
- ▶ heuristics: use the equivalence rules to push down selections and projections, and delay cartesian products
- ▶ constraints: left-deep plans, nested-loops and sort-merge join only
  - ▶ left-deep plans facilitate pipelining of output of each operator into the next operator, without materializing intermediate result

# Query optimization in System R

$$A \bowtie B \bowtie C \bowtie D$$


# Query optimization in System R

$$A \bowtie B \bowtie C \bowtie D$$






# Query optimization in System R

$A \bowtie B \bowtie C \bowtie D$




# Query optimization in System R

$A \bowtie B \bowtie C \bowtie D$





# Query optimization in System R

$$A \bowtie B \bowtie C \bowtie D$$




bushy



# Query optimization in System R

$$A \bowtie B \bowtie C \bowtie D$$



left-deep



# Query optimization in System R

- ▶ left-deep plans differ only in the order of relations, the access method (file-scan or index) for each relation, and the join method for each join

# Query optimization in System R

- ▶ left-deep plans differ only in the order of relations, the access method (file-scan or index) for each relation, and the join method for each join
- ▶ in spite of pruning, the resulting search space is still exponential in the number of relations (as we saw earlier)
  - ▶ actually is  $\mathcal{O}(n2^n)$  in practice
  - ▶ with  $n = 10$ , this is roughly 10,000 plans (compared with 59,000 and 17,600,000,000)

# Query optimization in System R

enumeration of left-deep plans: basic idea

- ▶ identify cheapest way to access each single relation in the query

# Query optimization in System R

enumeration of left-deep plans: basic idea

- ▶ identify cheapest way to access each single relation in the query
- ▶ for every access method and join predicate, find the cheapest way to join in a second relation

# Query optimization in System R

enumeration of left-deep plans: basic idea

- ▶ identify cheapest way to access each single relation in the query
- ▶ for every access method and join predicate, find the cheapest way to join in a second relation
- ▶ ...

# Query optimization in System R

enumeration of left-deep plans, for  $N$  relations

- ▶ Pass 1: find cheapest 1-relation plan for each relation
- ▶ Pass 2: find cheapest way to join result of each 1-relation plan (as outer) to another relation (output: all 2-relation plans)
- ▶ ...
- ▶ Pass  $N$ : find cheapest way to join result of a  $(N - 1)$ -relation plan (as outer) to the  $N$ th relation (output: all  $N$ -relation plans)

for each subset of relations, retain only cheapest plan overall, using  $\mathcal{O}(2^N)$  space

# Query optimization in System R

Suppose in our running example that we have 102 buffer pages available, and

- ▶ The Emp table has 10,000 tuples, on 1000 pages, with a clustered B-tree index on EID.
- ▶ The Works table has 20,000 tuples, on 2000 pages, sorted on  $\langle \text{EID}, \text{CID} \rangle$ .
- ▶ The Company table has 10,000 tuples, on 1000 pages, with a clustered B-tree index on CID.

# Query optimization in System R

Pass 1: finding cheapest 1-relation plans

- ▶ for Emp, we have the local predicate  $e.city = \text{'Best'}$ . However, with only an index on EID, the cheapest plan is a file-scan (with on-the fly application of the selection predicate



# Query optimization in System R

Pass 1: finding cheapest 1-relation plans

- ▶ for Emp, we have the local predicate  $e.city = \text{'Best'}$ . However, with only an index on EID, the cheapest plan is a file-scan (with on-the fly application of the selection predicate)
- ▶ likewise, we have file-scans as cheapest access method for Works, and Company

# Query optimization in System R

Pass 2: finding cheapest 2-relation plans

- ▶ we don't consider  $\{E, C\}$  since this is just a cartesian product

# Query optimization in System R

Pass 2: finding cheapest 2-relation plans

- ▶ we don't consider  $\{E, C\}$  since this is just a cartesian product
- ▶ for EW (or WE), we can perform merge-join at cost  $1000 + 2000 = 3000$
- ▶ for WE, we can perform index nested-loop-join at cost  $2000 + 3 \cdot 20,000 = 62,000$
- ▶ for EW, we can perform block-nested-loops-join at cost  $1000 + 2000 \cdot (1000/100) = 21,000$

# Query optimization in System R

Pass 2: finding cheapest 2-relation plans

- ▶ for WC, we can perform index nested-loop-join at cost  $2000 + 3 \cdot 20,000 = 62,000$
- ▶ for WC (or CW), we can perform sort-merge-join at cost  $4000 + 2000 + 1000 = 7000$
- ▶ for CW, we can perform block-nested-loops-join at cost  $1000 + 2000 \cdot 10 = 21,000$

# Query optimization in System R

Pass 2: finding cheapest 2-relation plans

- ▶ for WC, we can perform index nested-loop-join at cost  $2000 + 3 \cdot 20,000 = 62,000$
- ▶ for WC (or CW), we can perform sort-merge-join at cost  $4000 + 2000 + 1000 = 7000$
- ▶ for CW, we can perform block-nested-loops-join at cost  $1000 + 2000 \cdot 10 = 21,000$

So, best cost for  $\{E, W\}$  is 3000, and for  $\{W, C\}$  is 7000

# Query optimization in System R

Pass 3: finding cheapest 3-relation plans

- ▶ for EW joined with C, assuming 2000 pages in result of EW, we can perform sort merge-join at cost  $2000 + 4000 + 1000 = 7000$

# Query optimization in System R

Pass 3: finding cheapest 3-relation plans

- ▶ for EW joined with C, assuming 2000 pages in result of EW, we can perform sort merge-join at cost  $2000 + 4000 + 1000 = 7000$
- ▶ for WC joined with E, assuming 2000 pages in result of WC, we can perform sort merge-join at cost  $2000 + 4000 + 1000 = 7000$

# The life of a query: logical optimization

So, best overall plan is

- ▶ a merge-join of Emp and Works (with local predicates applied)
- ▶ followed by a sort-merge-join with Company,
- ▶ followed by the final projection

at cost  $3000 + 7000 = 10000$  I/Os



# Summary

- ▶ Query optimizer is at the heart of the query engine
- ▶ the paradigm typically followed is cost-based optimization
- ▶ System R follows a dynamic programming approach to plan space search
  - ▶ other practical approaches include randomized (e.g., simulated annealing) and genetic algorithms

# Summary

- ▶ Query optimizer is at the heart of the query engine
- ▶ the paradigm typically followed is cost-based optimization
- ▶ System R follows a dynamic programming approach to plan space search
  - ▶ other practical approaches include randomized (e.g., simulated annealing) and genetic algorithms

*Next time (Friday 22 May):* distributed data management